

John Sandsjö

DE24 - Datamodellering

Labb - databas för yrkeshögskolan YrkesCo

Agenda

Background

Business objective

Business requirements

Modelling

- Conceptual model

- Logical model

- Physical model

Implementation

- Creating tables

- Adding constraints

The data

Normalisation

Some query results and why they matter to the business

Note

The information in this presentation can also be read in the [yh_labbar_report.md](#) file in my GitHub repo

Preview Code Blame 115 lines (92 loc) · 6.11 KB Code 55% faster with GitHub Copilot Raw Copy Download Edit More

Build a database for YrkesCo


This document describes the process of creating a database model for the school YrkesCo.

Table of Contents

- [1. Business requirements](#)
- [2. Conceptual model](#)
- [3. Relationships statements for each entity](#)
- [4. Logical model](#)
- [5. Physical model](#)
- [6. Arguing for normalisation](#)
- [7. Creating database](#)

Business requirements

- om studenter, förnamn, efternamn, personnummer, email
- utbildare kan vara konsulter
- de planerar att anställa fasta utbildare (BONUS)
- utbildningsledare och deras personuppgifter
- utbildningsledare har hand om 3 klasser
- kurser med namn, kurskod, antal poäng, kort beskrivning av kursen
- program har ett antal kurser knutna till sig

 Prettier

Background

YrkesCo is a Swedish school with focus on Data and AI. It currently has two locations, one in Gothenburg and one in Stockholm, with possibly more sites to be open in the future. The schools program are location specific but the school also offer standalone courses to the students that can be taken from both sites.



Business objectives

Build a scalable database for YrkesCo. The model should adhere to future business needs like being able to scaling the operations to more sites as well as adhering to Personal Identifiable Information (PII).

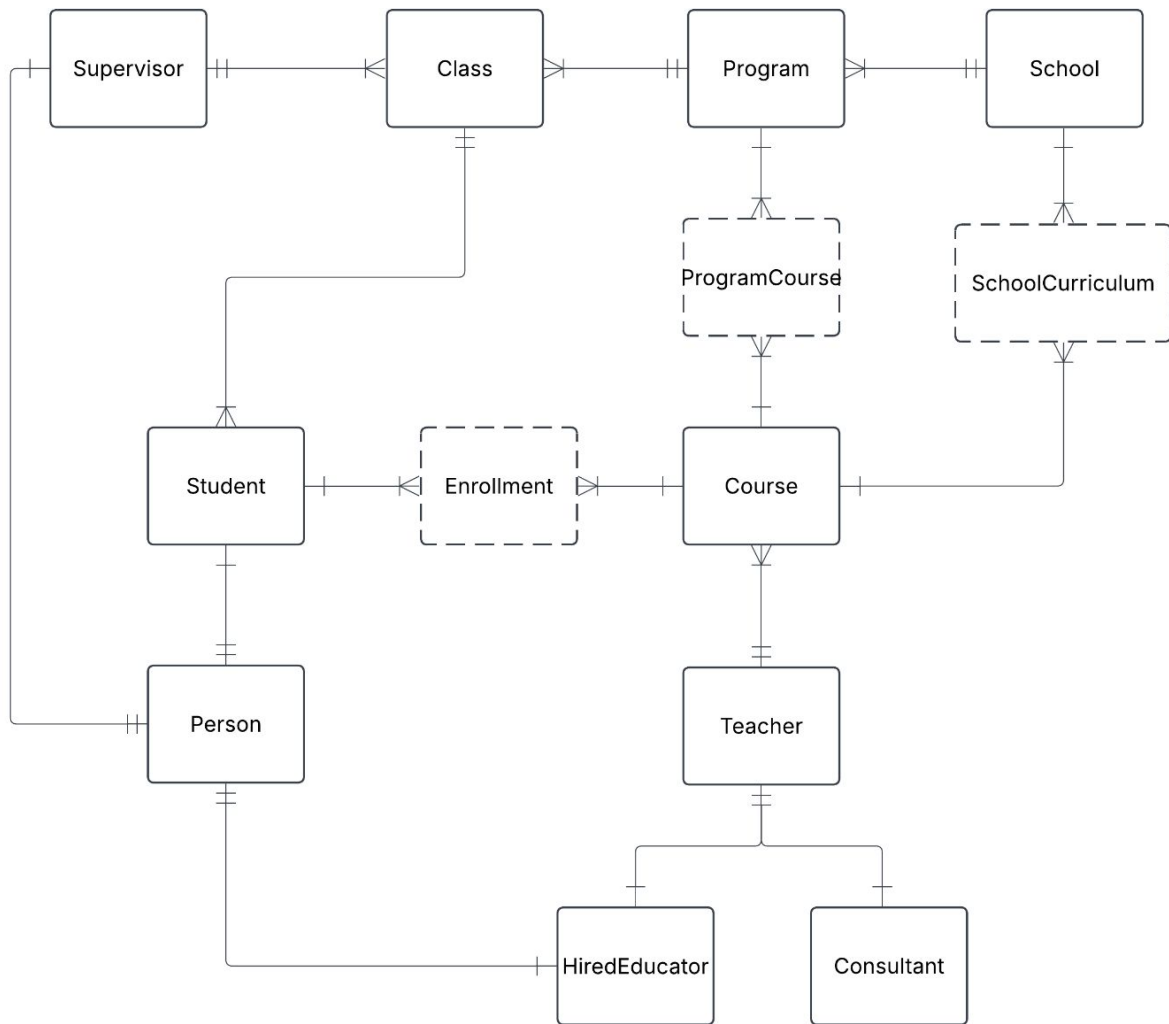

Business requirements, provided in task

- om studenter, förnamn, efternamn, personnummer, email
- utbildare kan vara konsulter
- de planerar att anställa fasta utbildare (BONUS)
- utbildningsledare och deras personuppgifter
- utbildningsledare har hand om 3 klasser
- kurser med namn, kurskod, antal poäng, kort beskrivning av kursen
- program har ett antal kurser knutna till sig
- ett program blir beviljat i tre omgångar, dvs att det finns 3 klasser
- det finns även fristående kurser (BONUS)
- konsulter, deras företag, företagsinfo som organisationsnummer, har F-skatt address, hur mycket de tar i arvode per timma
- YrkesCo har två anläggningar, en i göteborg och en i stockholm, i framtiden kanske de kommer expandera till flera orter (BONUS)

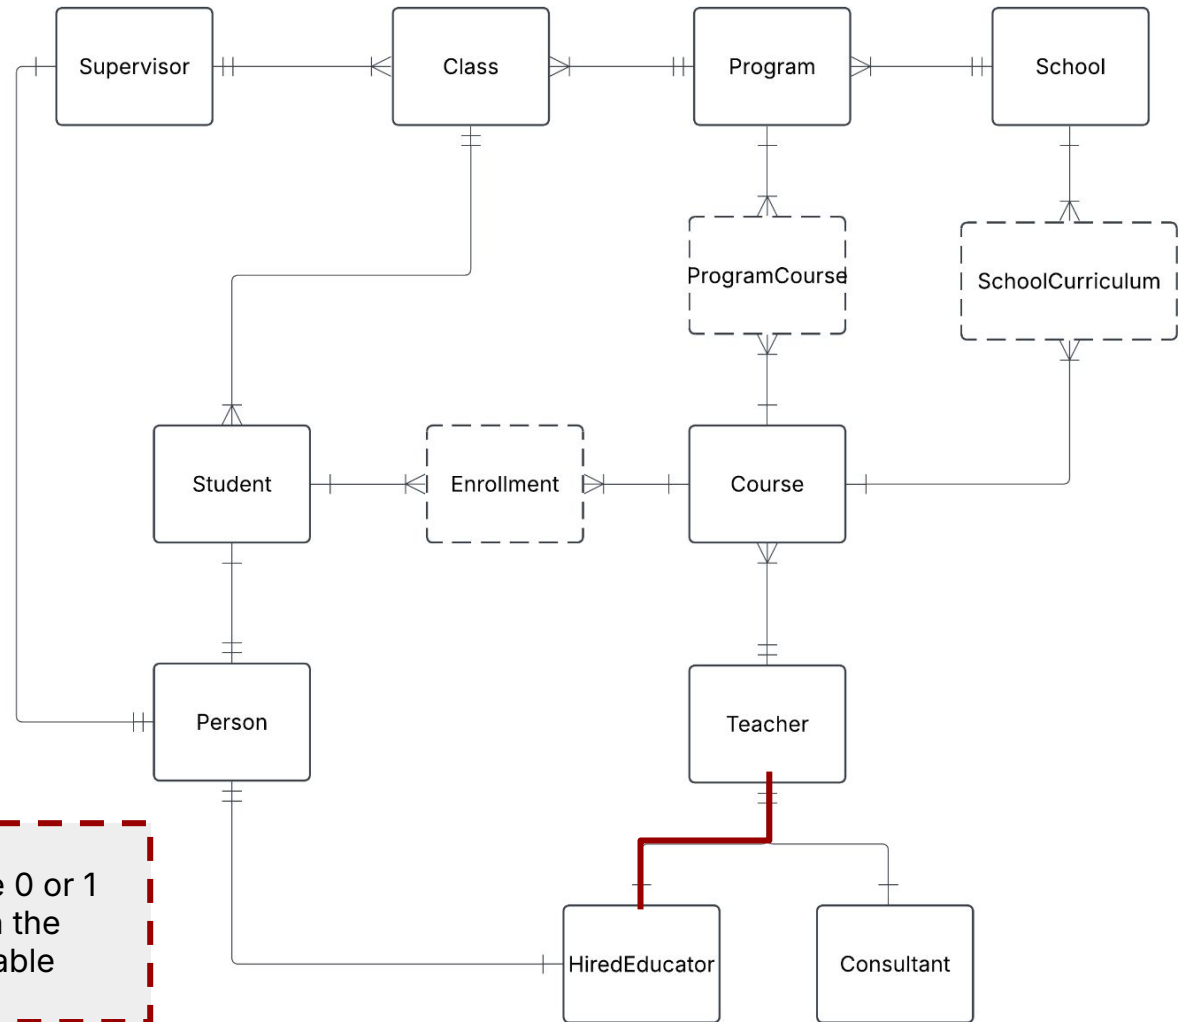

Additional business requirements

- A course can have only one teacher
- Standalone course can be taken online by students in any of the schools.
But courses are only open for students enrolled in the school.
- A program belongs to one school site
- A course can belong to many programs (e.g. Python fundamentals)

Conceptual model

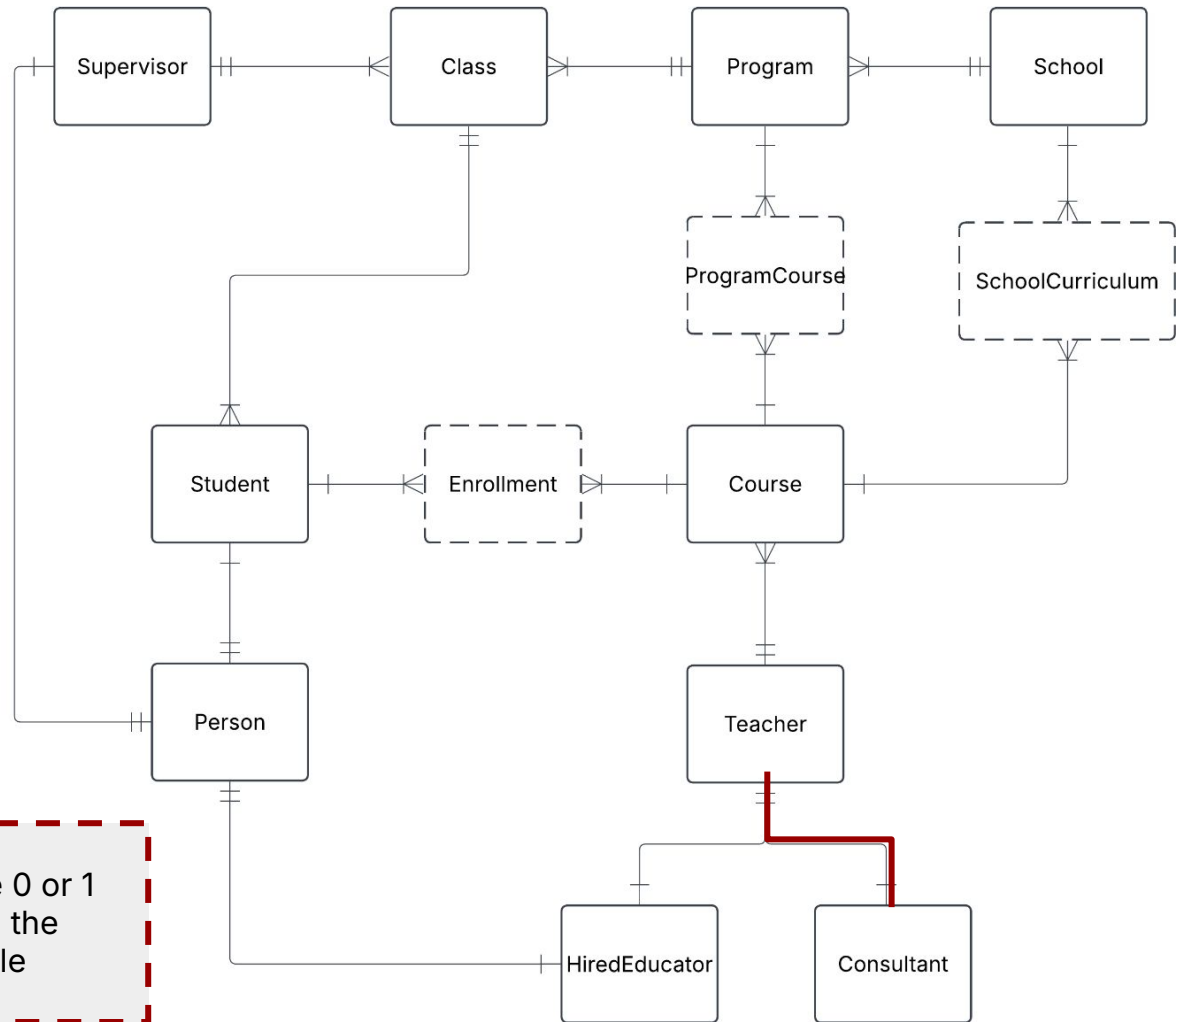



Conceptual model



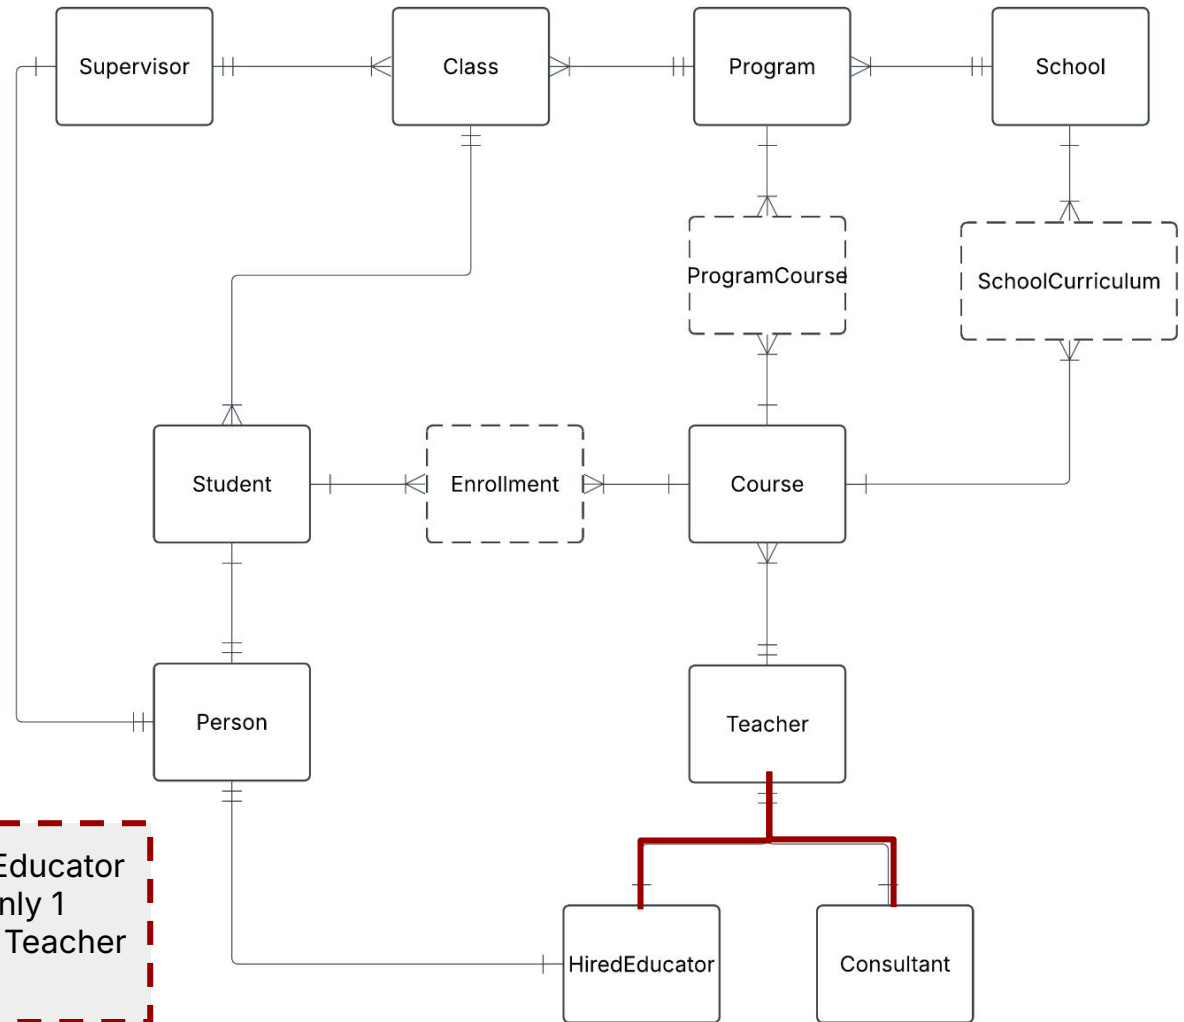

A teacher can have 0 or 1 representation in the HiredEducator table

Conceptual model



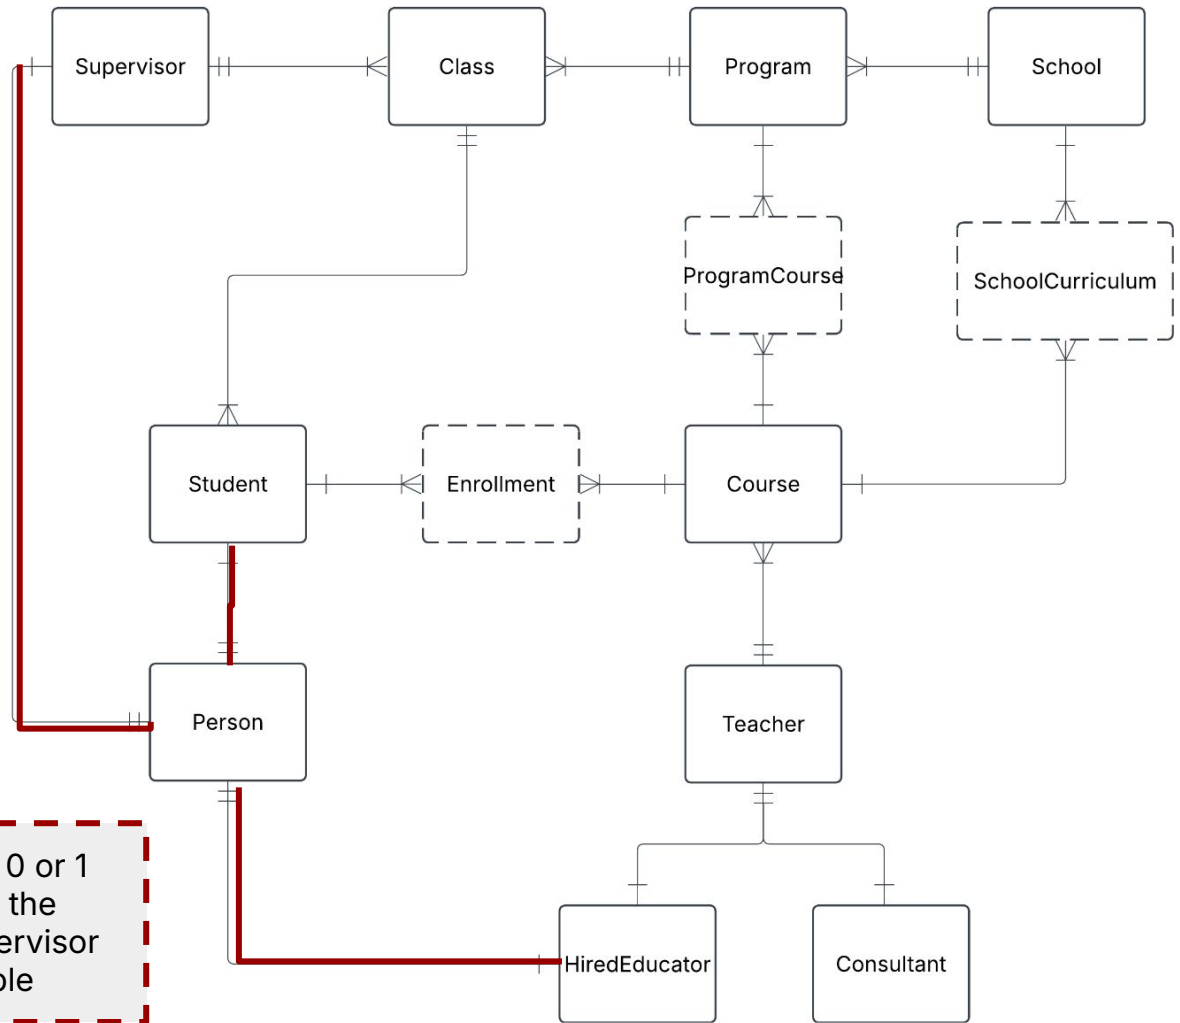

A teacher can have 0 or 1 representation in the Consultant table

Conceptual model



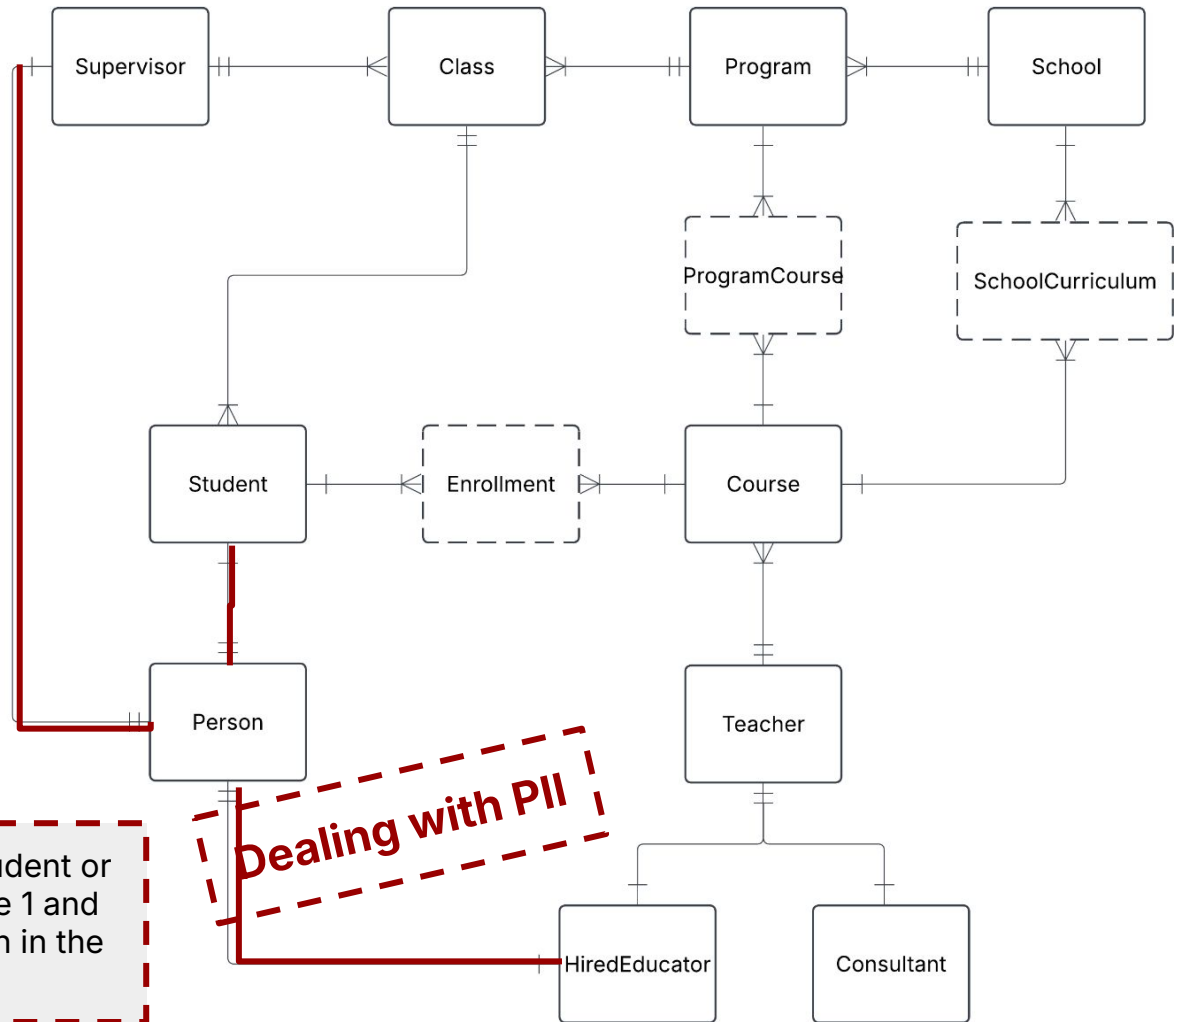

A Consultant / HiredEducator can have 1 and only 1 representation in the Teacher table

Conceptual model



A Person can have 0 or 1 representation in the HiredEducator, Supervisor and Student table

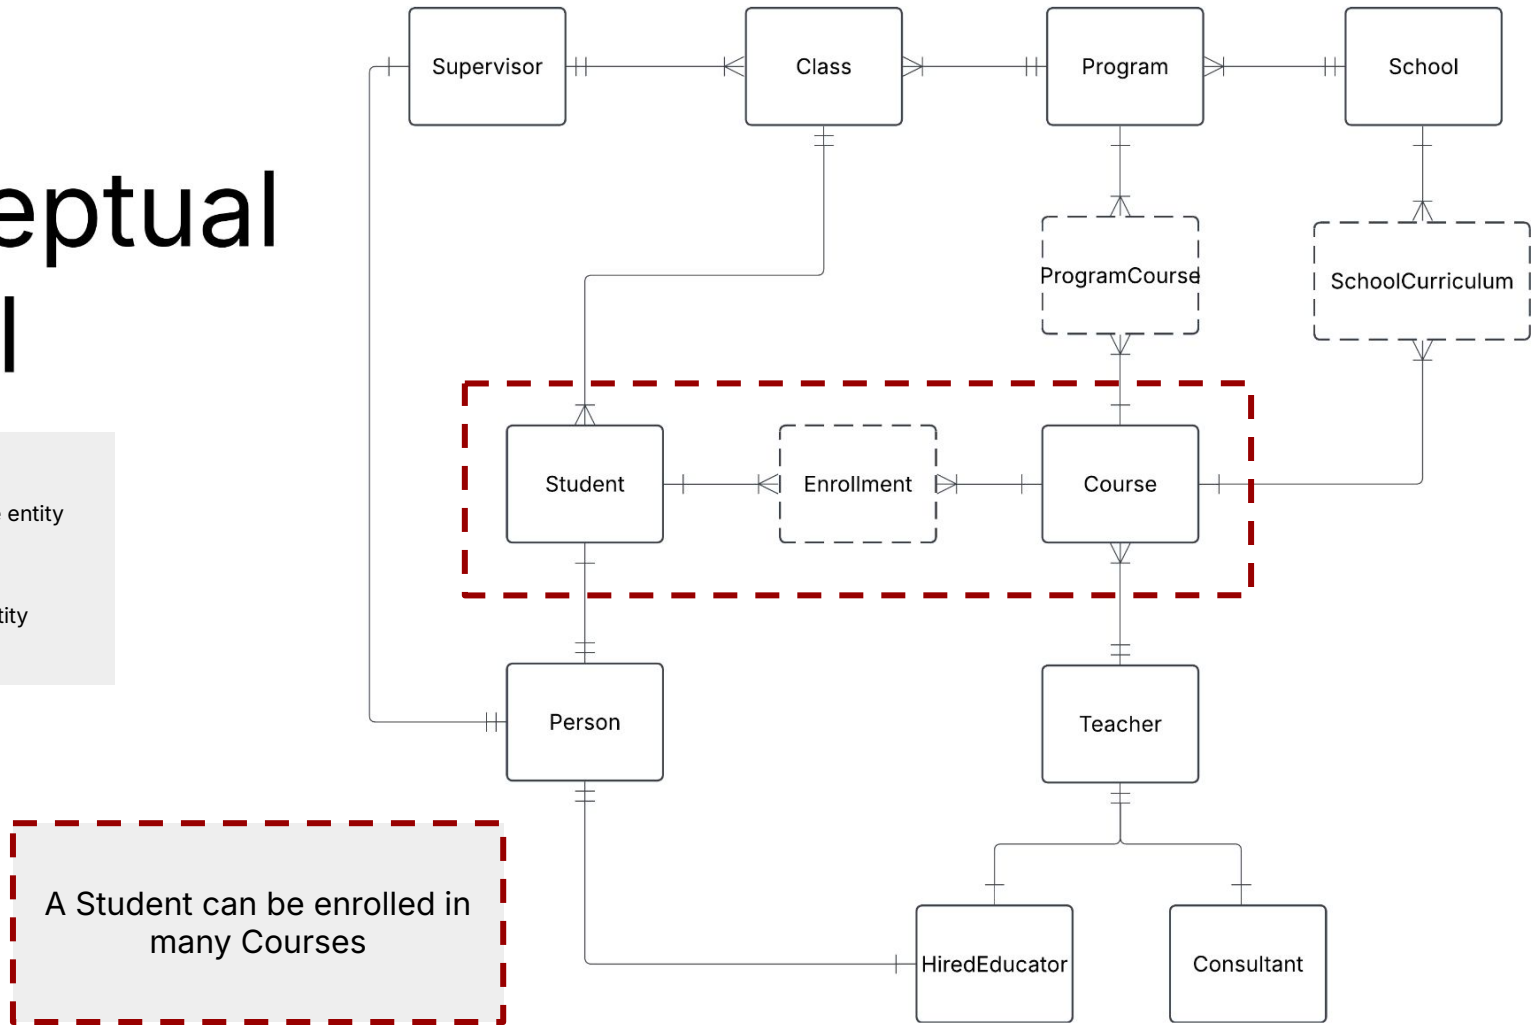

Conceptual model



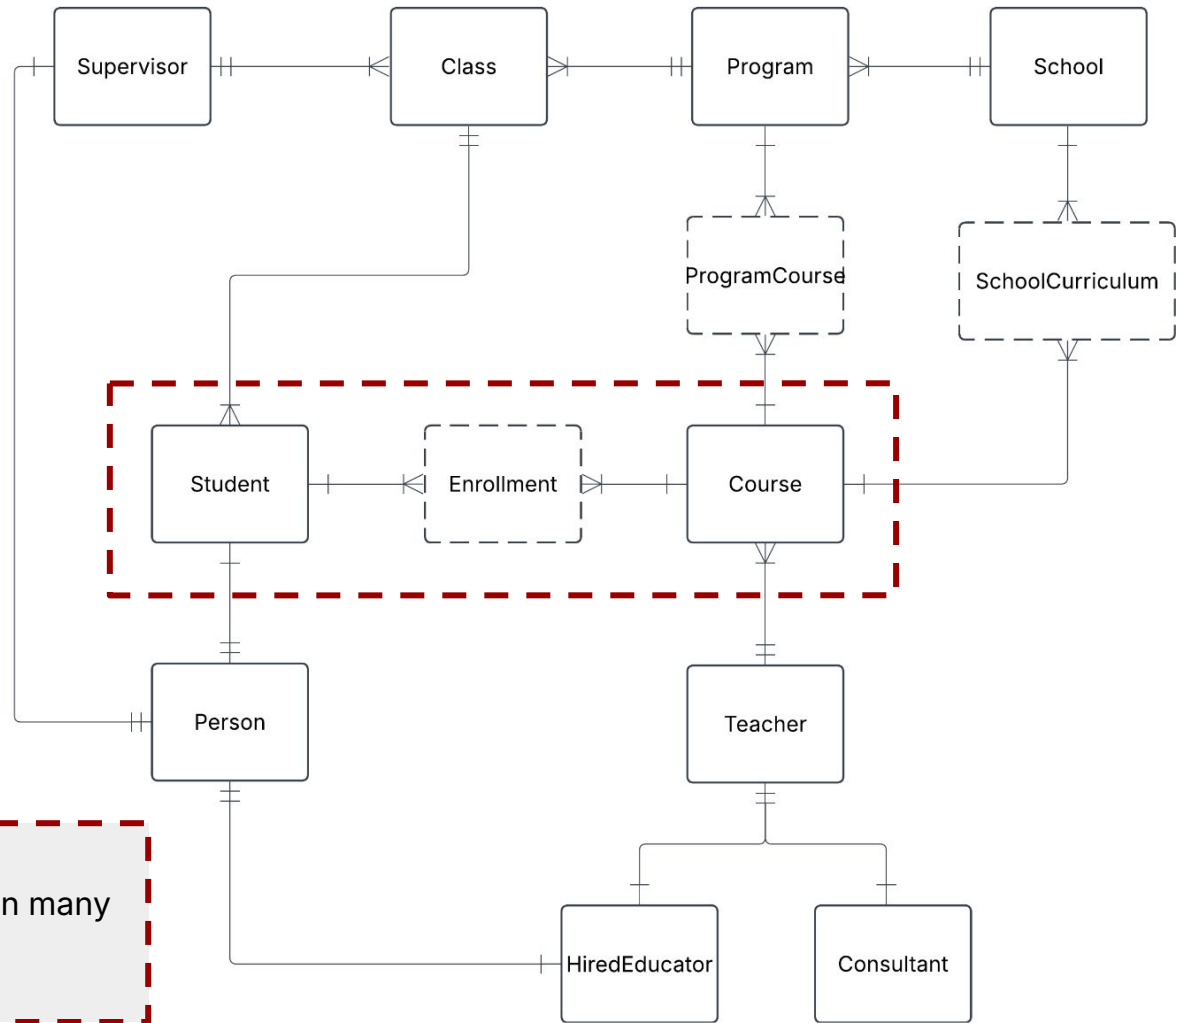

A HiredEducator, Student or Supervisor can have 1 and only 1 representation in the Person table

Dealing with PII

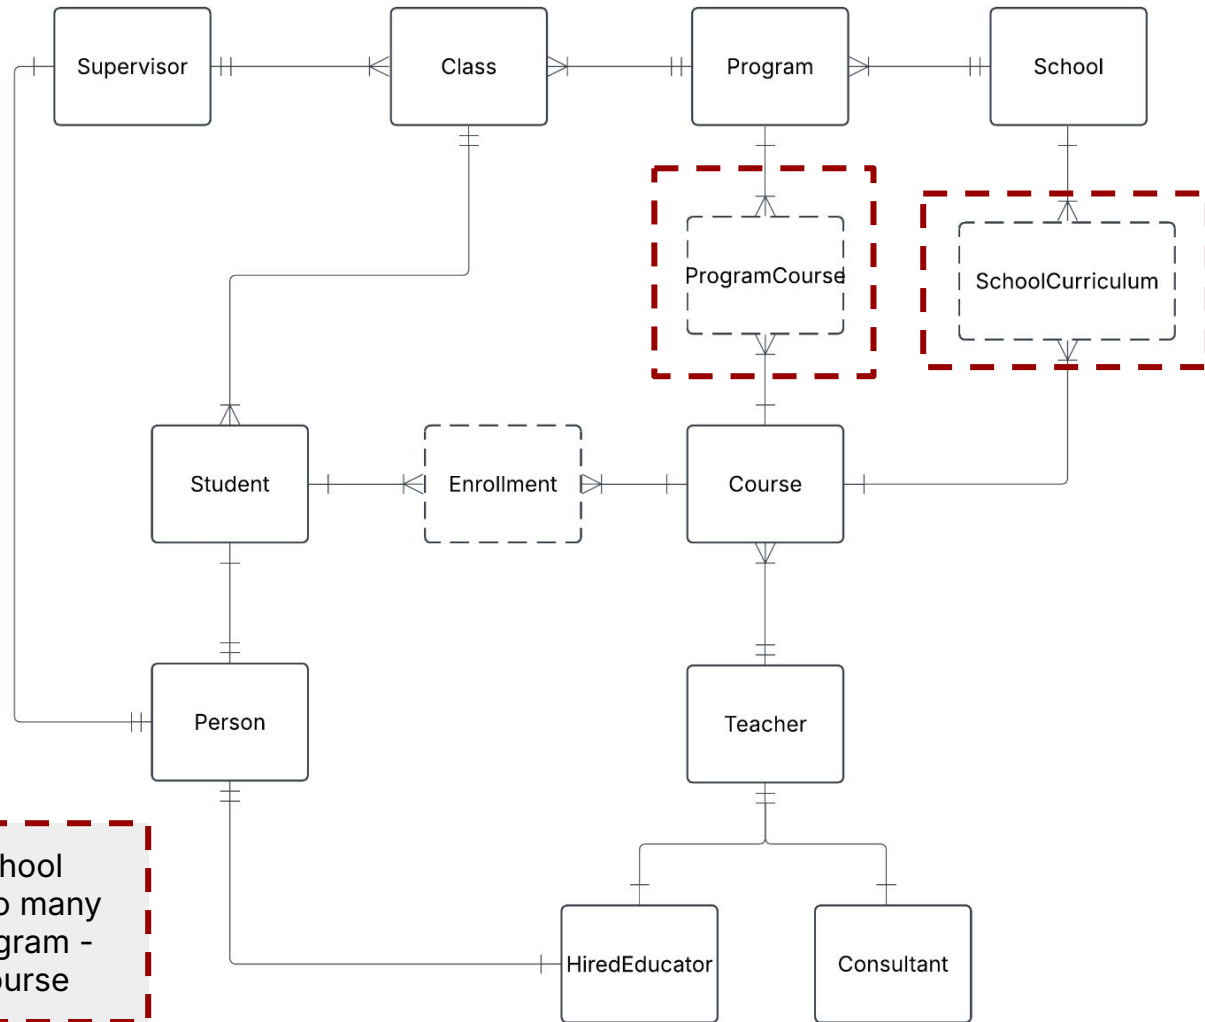

Conceptual model



Conceptual model

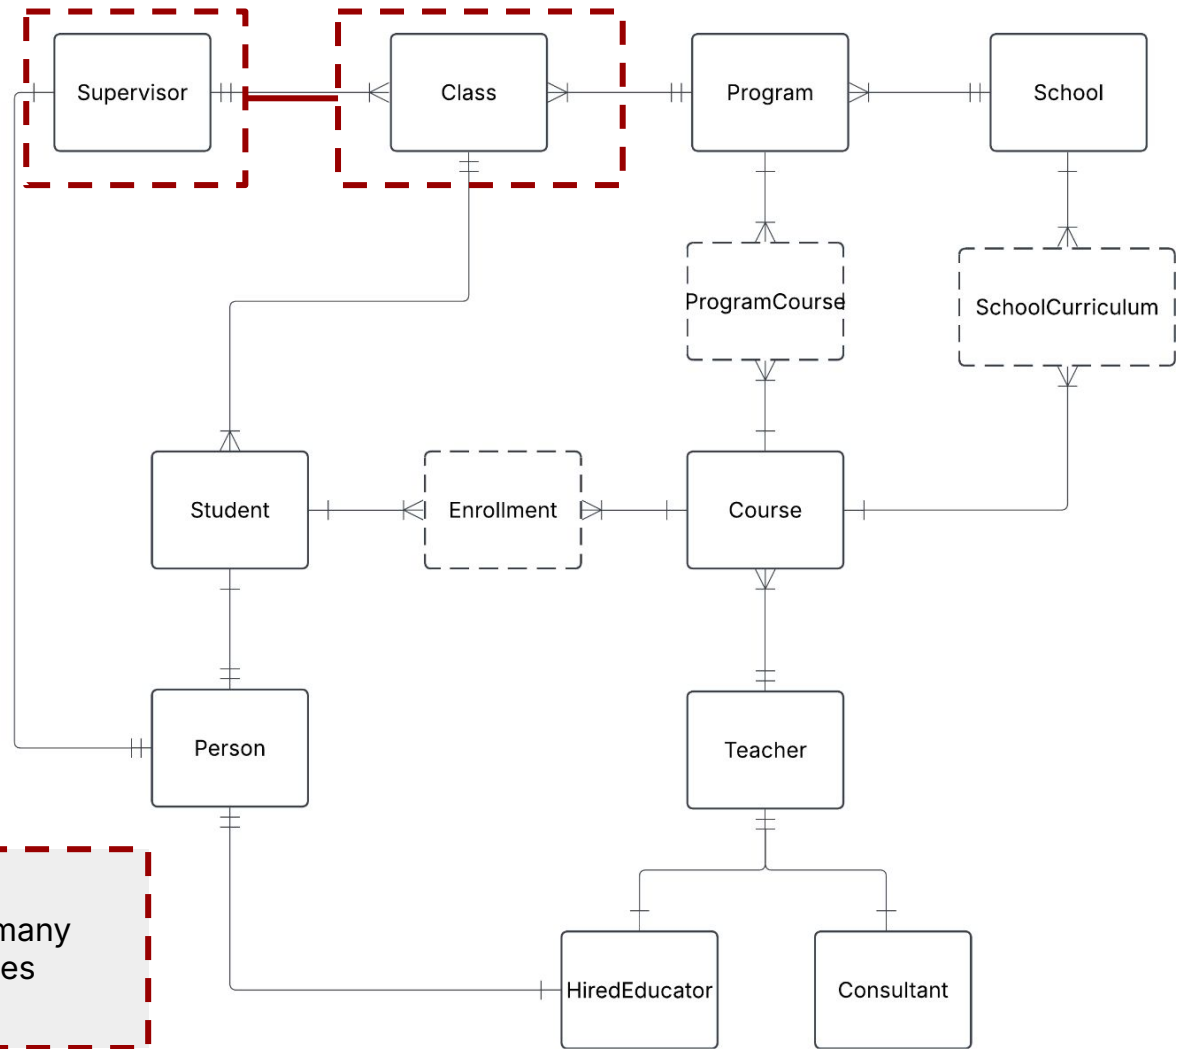



Conceptual model



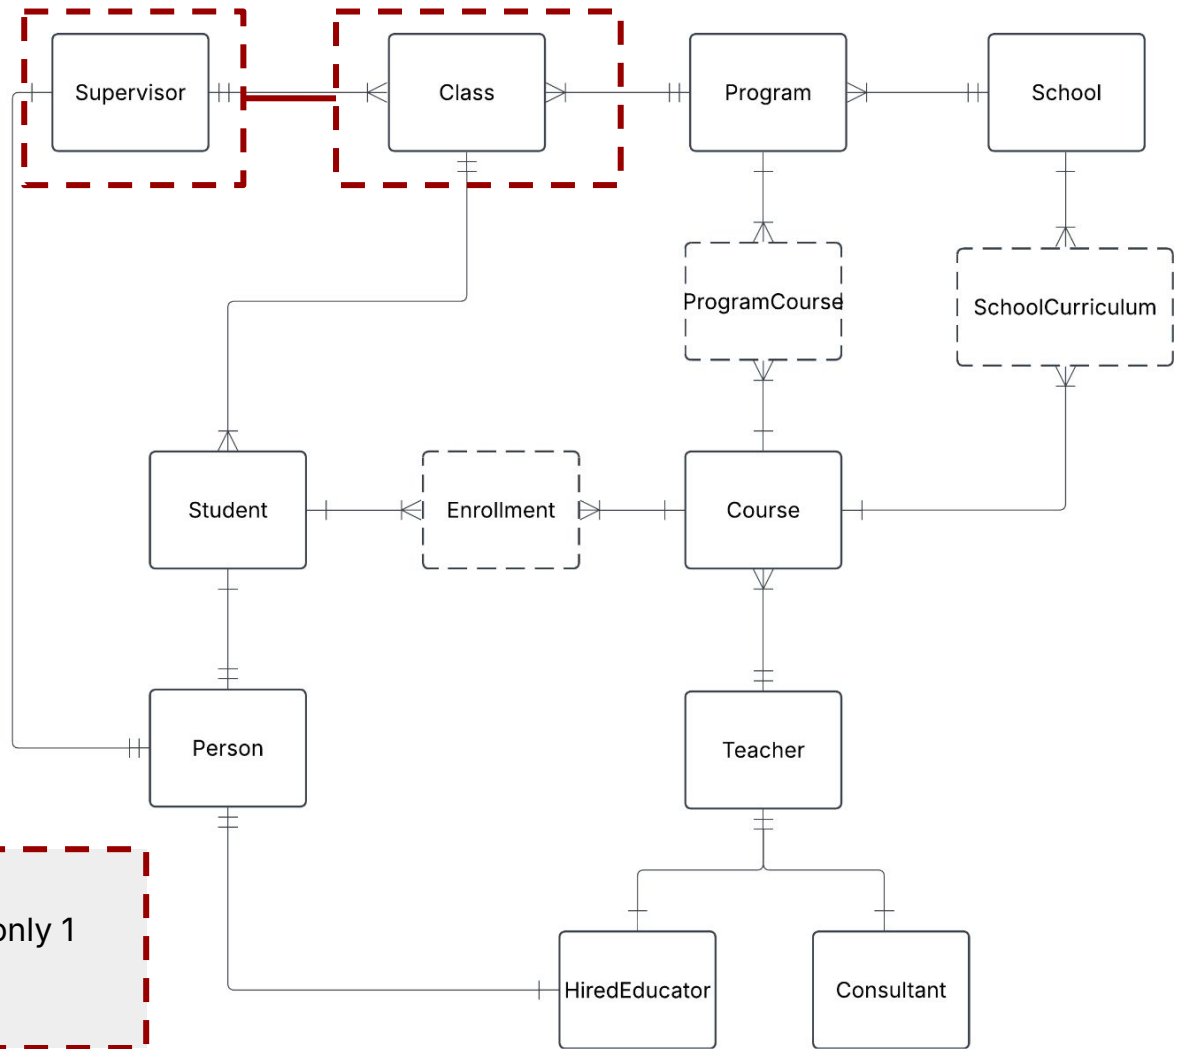

ProgramCourse and School Curriculum solves many to many relationship between Program - Course and School - Course

Conceptual model




A Supervisor can have many
(maximum of 3) classes

Conceptual model

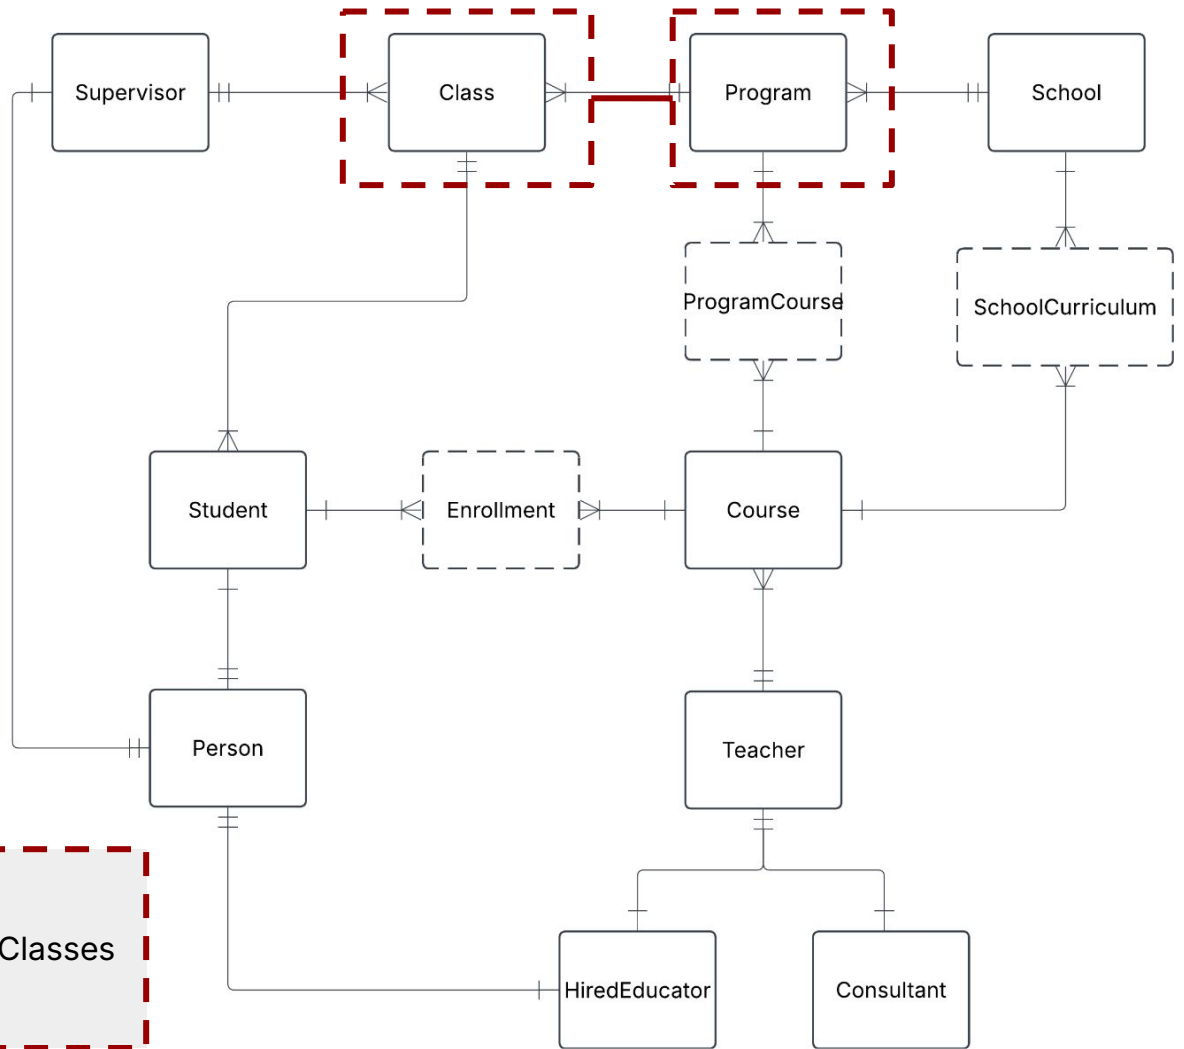


A Class can have 1 and only 1 Supervisor

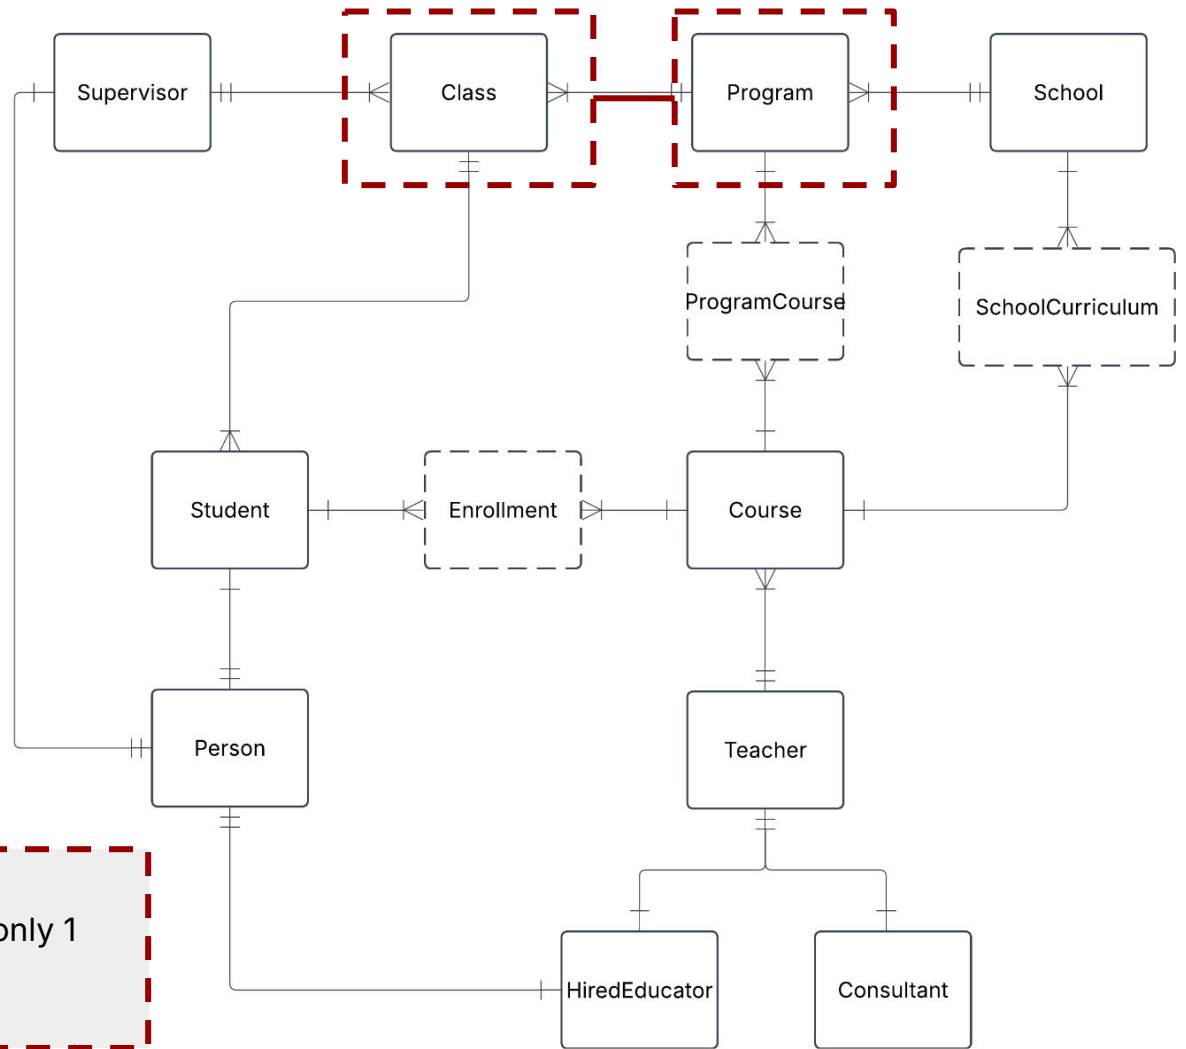

Conceptual model




A Program can have many Classes




Conceptual model




Conceptual model




A Program can be held in 1 and only 1 School



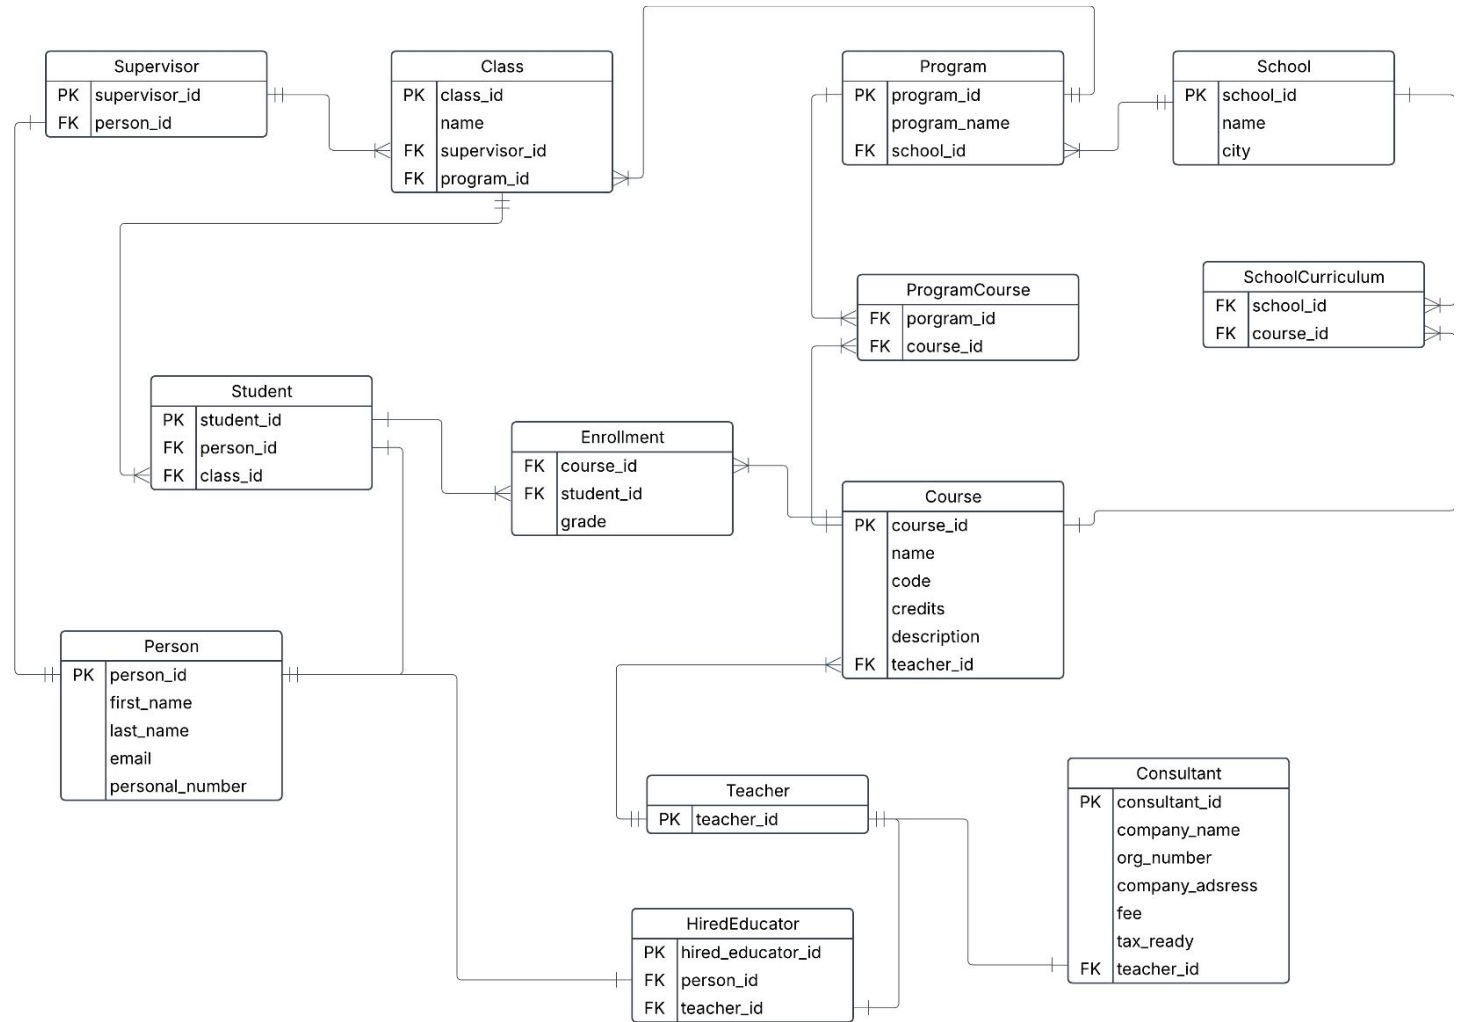
Conceptual model



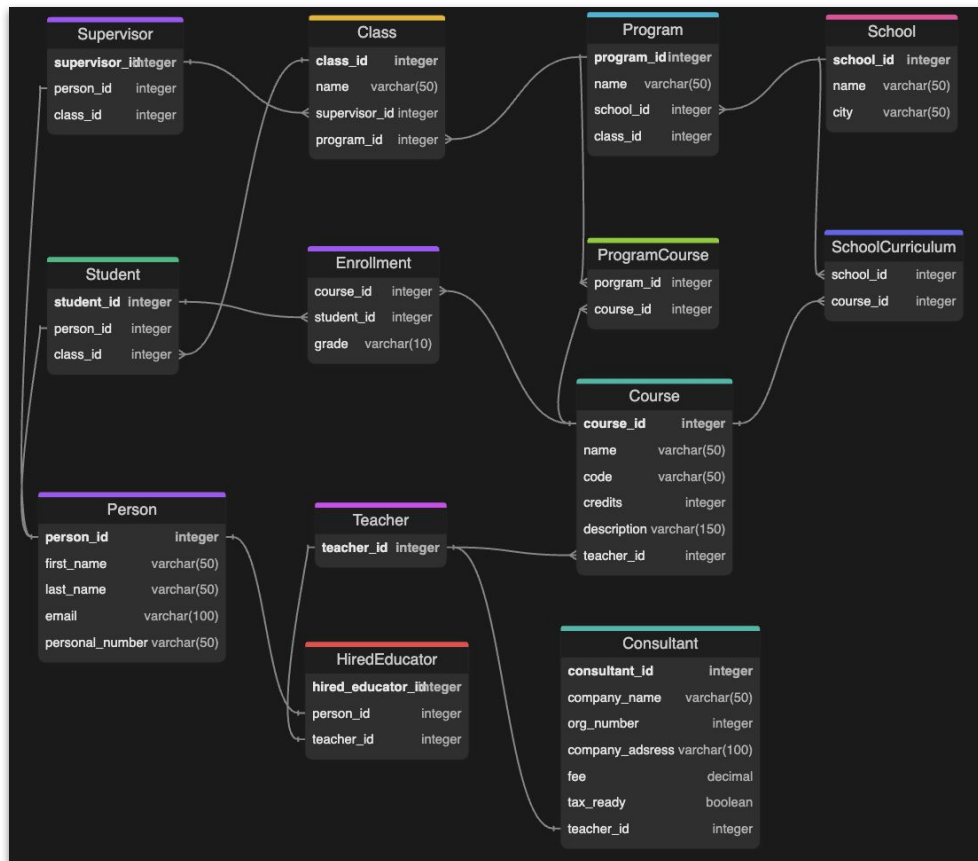
A School can have many Programs



Logical model



Physical model



Physical model

```
TABLE Supervisor {
  supervisor_id integer [PRIMARY KEY]
  person_id integer [ref: - Person.person_id]
  class_id integer
}

TABLE School {
  school_id integer [PRIMARY KEY]
  name varchar(50)
  city varchar(50)
}

TABLE Program {
  program_id integer [PRIMARY KEY]
  name varchar(50)
  school_id integer [ref: > School.school_id]
  class_id integer
}

TABLE Class {
  class_id integer [PRIMARY KEY]
  name varchar(50)
  supervisor_id integer [ref: > Supervisor.supervisor_id]
  program_id integer [ref: > Program.program_id]
}
```

Implementation

Creating tables

```
1 CREATE SCHEMA IF NOT EXISTS yh_labbb;
2
3 SET search_path TO yh_labbb;
4
5 CREATE TABLE IF NOT EXISTS School (
6     school_id integer PRIMARY KEY,
7     school_name varchar(50) NOT NULL,
8     city varchar(50) NOT NULL
9 );
10
11 CREATE TABLE IF NOT EXISTS Teacher (
12     teacher_id integer PRIMARY KEY
13 );
14
15 CREATE TABLE IF NOT EXISTS Program (
16     program_id integer PRIMARY KEY,
17     program_name varchar(50) NOT NULL,
18     school_id integer,
19     FOREIGN KEY (school_id) REFERENCES School (school_id)
20 );
21
22 CREATE TABLE IF NOT EXISTS Course (
23     course_id integer PRIMARY KEY,
24     course_name varchar(50) NOT NULL,
25     code varchar(50),
26     credits integer,
27     description varchar(150),
28     teacher_id integer,
29     FOREIGN KEY (teacher_id) REFERENCES Teacher (teacher_id)
30 );
```

Implementation

Adding constraints

Trigger

| Column | Table "yh_labbb.class" Type | Collation | Nullable | Default |
|---------------|--------------------------------|-----------|----------|---------|
| class_id | integer | | not null | |
| class_name | character varying(50) | | not null | |
| supervisor_id | integer | | | |
| program_id | integer | | | |

Indexes:
"class_pkey" PRIMARY KEY, btree (class_id)
Foreign-key constraints:
"class_program_id_fkey" FOREIGN KEY (program_id) REFERENCES yh_labbb.program(program_id)
"class_supervisor_id_fkey" FOREIGN KEY (supervisor_id) REFERENCES yh_labbb.supervisor(supervisor_id)
Referenced by:
TABLE "yh_labbb.student" CONSTRAINT "student_class_id_fkey" FOREIGN KEY (class_id) REFERENCES yh_labbb.class(class_id)
Triggers:
before_insert_class BEFORE INSERT OR UPDATE ON yh_labbb.class FOR EACH ROW EXECUTE FUNCTION check_supervisor_limit()

NOT NULL

```
CREATE TABLE IF NOT EXISTS Course (  
  course_id integer PRIMARY KEY,  
  course_name varchar(50) NOT NULL,  
  code varchar(50),  
  credits integer,  
  description varchar(150),  
  teacher_id integer,  
  FOREIGN KEY (teacher_id) REFERENCES Teacher (teacher_id)  
);
```

UNIQUE

```
-- Adding unique constraints to personal number  
ALTER TABLE yh_labbb.person  
ADD CONSTRAINT unique_personal_number  
UNIQUE (personal_number);
```

| | constraint_name name | constraint_type character varying |
|---|-------------------------|--------------------------------------|
| 1 | course_pkey | PRIMARY KEY |
| 2 | course_teacher_id_fkey | FOREIGN KEY |
| 3 | unique_course_code | UNIQUE |
| 4 | 17343_17565_1_not_null | CHECK |
| 5 | 17343_17565_2_not_null | CHECK |

The data

- Dummy data generated in LLM ingested with the purpose to test the relationships and joins
- Methodology for data ingestion

Prompt:

I am modelling a database for a school. See below my sql statement for creating my tables and its cardinalities. I want you to create dummy data in csv's for each of these tables? The dummy data should have the theme Data and AI and courses could be things like AI, Analytics and other development courses. The course names should have a fun twist to the course names. The School's name is YrkesCo and have one location in Stockholm, called "YrkesCo Liljeholmen" and one location in Gothenburg, called "YrkesCo Lindholmen". The student and teachers name should be Swedish sounding names, the first letter in the first_name should be the same as the first letter in the last_name.

The data

- Dummy data generated in LLM ingested with the purpose to test the relationships and joins
- Methodology for data ingestion

- Used Gemini to generate dummy data in csv's for each Table
- Added the csv's to my local machine
- Created a bind mount in docker compose with source folder from bullet above and target folder in mnt folder in container

```
-- Example Enrollment csv  
COPY enrollment FROM '/mnt/Enrollment.csv' WITH (FORMAT CSV, HEADER);
```

- Copied the data from the csv's to the database using the COPY command in psql

Normalisation - why do we care about it?

- Reduce data redundancy and improve data integrity
 - Changes can be done in place and integrity is ensured
- Prevent anomalies when updating, inserting
 - Ensures high data quality over the entire lifecycle
- Normalisation is not a silver bullet for OLAP and Data Warehousing
 - For analytics purposes in OLAP systems and data warehousing, denormalization is sometimes desired to optimize read performance for analytical queries for quicker data retrieval.

Normalisation - how it is achieved?

| Normal Form | Requirments | Argument |
|-------------|---|---|
| 1NF | <ul style="list-style-type: none">✓ all tables have primary key✓ No repeating groups✓ Uniform column data✓ Row order does not matter | <i>Going through each table, all of them has a primary key, the junction tables has it by combining its foreign keys. Each attribute is of one data type and includes no groupings. Row order does not matter. Thus, it adhere to first normal form</i> |
| 2NF | <ul style="list-style-type: none">✓ 1NF✓ Non-prime attributes must be functionally dependent on entire primary key and not just part of it | <i>Each table is 1NF. No attributes is functionally determined by other than the primary key. Thus, 2NF is reached</i> |
| 3NF | <ul style="list-style-type: none">✓ 2NF✓ Non-prime attributes depends on the key, the whole key and nothing but the key. | <i>Each table is 2NF. There are no transitive dependencies of the attributes. Thus, 3NF is reached</i> |

Some query results and why they matter to the business

Query

for each school, list the program names and the courses offered within those programs?

Why it matters

This is showing that the school sites are having distinct programs as modelled. It also shows that more school sites can be added in the future if the business grows.

| | school_name character varying | program_name character varying | course_name character varying | code character varying | description character varying |
|---|----------------------------------|-----------------------------------|----------------------------------|------------------------------|----------------------------------|
| 1 | YrkesCo Liljeholmen | Data Science Mastery | Pythonic Data Delights | PYDD101 | Python for data analysis |
| 2 | YrkesCo Liljeholmen | Data Science Mastery | Neural Network Nirvana | NNN201 | Deep learning foundations |
| 3 | YrkesCo Liljeholmen | AI Innovation Lab | Neural Network Nirvana | NNN201 | Deep learning foundations |
| 4 | YrkesCo Liljeholmen | AI Innovation Lab | AI & Algorithm Alchemy | AAAA501 | Advanced AI and algorithms |
| 5 | YrkesCo Lindholmen | Full Stack Web Dev | React & Roll | RR301 | Web development with React |
| 6 | YrkesCo Lindholmen | Data Analytics Pro | Data Dive & Discover | DDD401 | Data visualization & exploration |
| 7 | YrkesCo Liljeholmen | Machine Learning Wizardry | AI & Algorithm Alchemy | AAAA501 | Advanced AI and algorithms |
| 8 | YrkesCo Lindholmen | Cloud Engineering | Cloud Computing Chronicles | CCC601 | Cloud computing essentials |
| 9 | YrkesCo Liljeholmen | Data Science Mastery | Database Design Dazzle | DDDD701 | Database design and SQL |

Query

what courses belong to more than one program?

Why it matters

This is showing that the school can have shared courses across programs as defined in the business requirements.

| | course_name character varying | nr_courses bigint |
|----------|---|-----------------------------|
| 1 | AI & Algorithm Alchemy | 2 |
| 2 | Neural Network Nirvana | 2 |

Query

What courses belong to which program and which courses are standalone?

Why it matters

This is showing that the school can offer standalone courses not part of any program. Open for any student to attend.

| | course_name character varying | code character varying | program_name character varying |
|----|----------------------------------|---------------------------|-----------------------------------|
| 1 | Pythonic Data Delights | PYDD101 | Data Science Mastery |
| 2 | Neural Network Nirvana | NNN201 | Data Science Mastery |
| 3 | Neural Network Nirvana | NNN201 | AI Innovation Lab |
| 4 | AI & Algorithm Alchemy | AAAA501 | AI Innovation Lab |
| 5 | React & Roll | RR301 | Full Stack Web Dev |
| 6 | Data Dive & Discover | DDD401 | Data Analytics Pro |
| 7 | AI & Algorithm Alchemy | AAAA501 | Machine Learning Wizardry |
| 8 | Cloud Computing Chronicles | CCC601 | Cloud Engineering |
| 9 | Database Design Dazzle | DDDD701 | Data Science Mastery |
| 10 | Intro to RAGs | IR20 | <i>null</i> |

Query

For each program, list the names of the students and their personal number enrolled in the classes within that program.

Why it matters

This is showing that simplicity of joining tables to get lists of personal information per course

| | program_name character varying | course_name character varying | first_name character varying | last_name character varying | personal_number character varying |
|----|-----------------------------------|----------------------------------|------------------------------------|-----------------------------------|--------------------------------------|
| 1 | Data Science Mastery | Pythonic Data Delights | Niklas | Nilsson | 198402204567 |
| 2 | Data Science Mastery | Pythonic Data Delights | Greta | Gustafsson | 199107156789 |
| 3 | Data Science Mastery | Pythonic Data Delights | Fredrik | Forsberg | 198306102345 |
| 4 | Data Science Mastery | Neural Network Nirvana | Olivia | Olsson | 199703258901 |
| 5 | Data Science Mastery | Neural Network Nirvana | Ida | Isaksson | 199409254567 |
| 6 | Data Science Mastery | Neural Network Nirvana | Hanna | Holm | 198708200123 |
| 7 | AI Innovation Lab | Neural Network Nirvana | Olivia | Olsson | 199703258901 |
| 8 | AI Innovation Lab | Neural Network Nirvana | Ida | Isaksson | 199409254567 |
| 9 | AI Innovation Lab | Neural Network Nirvana | Hanna | Holm | 198708200123 |
| 10 | AI Innovation Lab | AI & Algorithm Alchemy | Fredrik | Forsberg | 198306102345 |
| 11 | AI Innovation Lab | AI & Algorithm Alchemy | Olivia | Olsson | 199703258901 |
| 12 | AI Innovation Lab | AI & Algorithm Alchemy | Niklas | Nilsson | 198402204567 |
| 13 | Full Stack Web Dev | React & Roll | Per | Persson | 198104302345 |
| 14 | Full Stack Web Dev | React & Roll | Karin | Karlsson | 199311052345 |
| 15 | Full Stack Web Dev | React & Roll | Johan | Jonsson | 198210308901 |
| 16 | Data Analytics Pro | Data Dive & Discover | Quirine | Qvist | 199805056789 |
| 17 | Data Analytics Pro | Data Dive & Discover | Maria | Magnusson | 199601150123 |
| 18 | Data Analytics Pro | Data Dive & Discover | Lars | Lindberg | 198612106789 |
| 19 | Machine Learning Wizardry | AI & Algorithm Alchemy | Fredrik | Forsberg | 198306102345 |
| 20 | Machine Learning Wizardry | AI & Algorithm Alchemy | Olivia | Olsson | 199703258901 |
| 21 | Machine Learning Wizardry | AI & Algorithm Alchemy | Niklas | Nilsson | 198402204567 |
| 22 | Cloud Engineering | Cloud Computing Chronicles | Greta | Gustafsson | 199107156789 |
| 23 | Cloud Engineering | Cloud Computing Chronicles | Quirine | Qvist | 199805056789 |
| 24 | Cloud Engineering | Cloud Computing Chronicles | Per | Persson | 198104302345 |
| 25 | Data Science Mastery | Database Design Dazzle | Hanna | Holm | 198708200123 |

Query

What supervisor belong to a specific class?

Why it matters

This is showing that a supervisor has 3 classes each. A trigger function is also added to enforce this business requirement.

| | class_name character varying | supervisor_id integer | first_name character varying | last_name character varying |
|---|---------------------------------|--------------------------|---------------------------------|--------------------------------|
| 1 | DS2023-A | 1 | Daniel | Dahlberg |
| 2 | AI2023-B | 1 | Daniel | Dahlberg |
| 3 | WD2023-C | 1 | Daniel | Dahlberg |
| 4 | AN2023-D | 2 | Robert | Rosén |
| 5 | ML2023-E | 2 | Robert | Rosén |
| 6 | CE2023-F | 2 | Robert | Rosén |

| | class_name character varying | program_name character varying | course_name character varying | credits integer | first_name character varying | last_name character varying | company_name character varying |
|---|---------------------------------|-----------------------------------|----------------------------------|--------------------|---------------------------------|--------------------------------|-----------------------------------|
| 1 | AI2023-B | AI Innovation Lab | Neural Network Nirvana | 40 | Birgitta | Bergman | <i>null</i> |
| 2 | AI2023-B | AI Innovation Lab | AI & Algorithm Alchemy | 30 | <i>null</i> | <i>null</i> | AIConsulting |

Query

What courses a specific class has and
what educator per course

Why it matters

This is showing that for a specific
program, the courses can be taught
by either a hired teacher or a
consultant.

Thank you for listening...